Statistics 341: Intro to Stat Computing and Exploratory Data Analysis with R

Lecture 2: R Data Structures, Functions and Packages

Brad Mcneney

2019-01-10

R objects

R objects

- In R, data structures and functions are all referred to as "objects".
- Objects are created with the assignment operator <−; e.g., x <- 1.
  - The objects a user creates from the R console are contained in the user's workspace, called the global environment.
  - Use ls() to see a list of all objects in the workspace.
  - Use rm(x) to remove object x from the workspace.
- We will discuss data structures first, and then functions and R packages.
- In between we will discuss logical and relational operators that can be used to subset data structures.

R Data Structures

R Data Structures

- ▶ Reference Wickham (2014), Advanced R, Chapter 2
- ▶ Focus on four common data structures: atomic vectors, lists, matrices and data frames.
- ▶ Atomic vectors and lists are 1d, while matrices and data frames are 2d objects
- ▶ R has no true scalars; e.g., in x<-1, x is a vector of length one.
- ▶ R also has an array data structure for higher dimensional elements that we will not discuss.
- ▶ Use str() to see the structure of an object

Types of objects

- ▶ All R objects have a "type", that describes how it is stored in computer memory.
- ▶ Common types we will encounter are "logical", "integer", "double", "character" and "list".
  - ▶ Find the type of an object with typeof().

```r
x <- 6 # stores as double by default
typeof(x)
```

```
## [1] "double"
```

```r
y <- 6L # The "L" suffix forces storage as integer
typeof(y)
```

```
## [1] "integer"
```

Type *versus* Mode

- ▶ In addition to the type of an object, there is its "mode".
- ▶ The mode of an object is generally the same as its type, but the modes are coarser.
    - ▶ For example, integer and double types are both of mode "numeric".
- ▶ I don't understand the need for mode.
    - ▶ The only reason I mention it is that the `str()` function sometimes reports the mode of an object, rather than its type, so we will frequently see reports of numeric objects.

```
mode(x)
```

```
## [1] "numeric"
```

```
mode(y)
```

```
## [1] "numeric"
```

Vectors

Vectors

- ▶ Vectors can be either atomic or list
    - ▶ The elements of an atomic vector must be the same type.
    - ▶ Lists can be comprised of multiple data types
- ▶ Empty vectors can be created by the vector() function:

```
# help("vector")
avec <- vector(mode="numeric",length=4)
lvec <- vector(mode="list",length=4)
```

- ▶ Data vectors can be created with c() or list():

```
avec <- c(52,200,77)
lvec <- list(52,200,77,c("grey","thin"))
```

Combining vectors

- ▸ Use c() to combine vectors

```r
c(avec,c(100,101))
```

```
## [1]  52 200  77 100 101
```

```r
c(lvec,TRUE)
```

```
## [[1]]
## [1] 52
##
## [[2]]
## [1] 200
##
## [[3]]
## [1] 77
##
## [[4]]
## [1] "grey" "thin"
##
## [[5]]
## [1] TRUE
```

Vector attributes

- Vectors have a type and length and, optionally, attributes such as names.
    - As we have seen, we find the type of an object with typeof().
    - Find the length of a vector with length().

```
typeof(avec)
```

```
## [1] "double"
```

```
length(avec)
```

```
## [1] 3
```

```
str(avec)
```

```
##  num [1:3] 52 200 77
```

```
typeof(lvec)
```

```
## [1] "list"
```

```
length(lvec)
```

```
## [1] 4
```

```
names(lvec) = c("age","weight","height","hair")
str(lvec)
```

```
## List of 4
##  $ age   : num 52
##  $ weight: num 200
##  $ height: num 77
##  $ hair  : chr [1:2] "grey" "thin"
```

▶ We can specify element names when creating a vector; e.g.:

```
lvec <- list(age=50,weight=200,height=77,hair=c("grey","thin"))
```

Factors

- The statistical concept of a factor is important in experimental design.
- Factors are implemented in R as atomic vectors with attributes class and levels:

```
trt <- factor(c("drug1","placebo","placebo","drug2"))
attributes(trt)
```

```
## $levels
## [1] "drug1"   "drug2"   "placebo"
##
## $class
## [1] "factor"
```

```
str(trt)
```

```
##  Factor w/ 3 levels "drug1","drug2",..: 1 3 3 2
```

- ▶ The levels are coded numerically (1, 2 and 3) with assigned labels ordered alphabetically ("drug1", "drug2" and "placebo") by default.
- ▶ You can specify an order to the factors with the `level` argument:

```
trt <- factor(c("drug1","placebo","placebo","drug2"),
              levels=c("placebo","drug1","drug2"))
trt
```

```
## [1] drug1   placebo placebo drug2
## Levels: placebo drug1 drug2
```

More on object class

- ▶ You can create your own `class` for an object.
- ▶ Such "meta-data" can be used to tell R how to handle the object; e.g., how to print it, summarize it, etc.

```
class(lvec) <- "prof"
lvec
```

```
## $age
## [1] 50
##
## $weight
## [1] 200
##
## $height
## [1] 77
##
## $hair
## [1] "grey" "thin"
##
## attr(,"class")
## [1] "prof"
```

Subsetting vectors and extracting elements

▶ Subset with [or by name:

```r
lvec[c(1,3)] # same as lvec[c("age","height")]
```

```
## $age
## [1] 50
##
## $height
## [1] 77
```

▶ Extract individual elements with [[, or $ for named objects:

```r
lvec[[4]]
```

```
## [1] "grey" "thin"
```

```r
lvec$hair
```

```
## [1] "grey" "thin"
```

Subsetting factors

- ▶ Subsetting may remove all instances of a level, but the level will be retained in the data structure

```
trt[1:3]
```

```
## [1] drug1   placebo placebo
## Levels: placebo drug1 drug2
```

- ▶ If subsetting is intended to remove a level of the factor, use drop=TRUE

```
trt[1:3,drop=TRUE]
```

```
## [1] drug1   placebo placebo
## Levels: placebo drug1
```

Subsetting and assignment

- You can combine subsetting and assignment to change the value of vectors

```
avec
```

```
## [1]  52 200  77
```

```
avec[2] <- 210
avec
```

```
## [1]  52 210  77
```

Assignment and lists

- To assign to a list element, use [[rather than [.
  - Assignment with [requires that the replacement element be of length 1; [[does not have this restriction

```
lvec[3:4] <- c("Hi","there")
lvec[3:4]
```

```
## $height
## [1] "Hi"
##
## $hair
## [1] "there"
```

```r
lvec[4] <- c("All","of","this")
```

```
## Warning in lvec[4] <- c("All", "of", "this"): number of items to replace is
## not a multiple of replacement length
```

```r
lvec[4] # Only used first element of replacement vector
```

```
## $hair
## [1] "All"
```

```r
lvec[[4]] <- c("All","of","this")
lvec[3:4]
```

```
## $height
## [1] "Hi"
##
## $hair
## [1] "All"   "of"    "this"
```

Coercion: atomic vectors to lists

- Atomic vectors can be coerced to lists with `as.list()`:

```
avec = c(age=50,weight=200,height=77)
avec
```

```
##    age weight height
##     50    200     77
```

```
as.list(avec)
```

```
## $age
## [1] 50
##
## $weight
## [1] 200
##
## $height
## [1] 77
```

Coercion: lists to atomic vectors

▶ Lists can be "flattened" into atomic vectors with unlist():

```
unlist(lvec)
```

```
##    age weight height  hair1  hair2  hair3
## "50"  "200"   "Hi"  "All"   "of" "this"
```

▶ Notice how the numeric values are coerced to the more flexible character type.
▶ The order of flexibility, from least to most, is logical, integer, numeric, character.

Coercion: factors to atomic vectors

- We saw how to use factor() to coerce an atomic vector to a factor.
- Use as.vector() to coerce a factor back to an atomic vector.
- The result is a character vector. You may need to use as.numeric() to coerce to numeric, if required.

```
a <-factor(c(2,1,1,2))
as.vector(a)
```

```
## [1] "2" "1" "1" "2"
```

```
as.numeric(as.vector(a))
```

```
## [1] 2 1 1 2
```

Matrices and data frames

Matrices and data frames

- Though both 2d objects, matrices and data frames are different enough that we will need to discuss them separately.
- The elements of a matrix must all be of the same type.
- Data frames are essentially lists where each list element has the same length. Thus data frames can include columns of varying type.

Matrices

▶ Matrices can be created with the matrix() function as in

```
A <- matrix(1:4,nrow=2,ncol=2)
A
```

```
##      [,1] [,2]
## [1,]    1    3
## [2,]    2    4
```

▶ Here 1:4 is the same as c(1,2,3,4)

▶ The default is to read the data vector into the matrix column-by-column. To read row-by-row instead use the `byrow=TRUE` argument:

```
A <- matrix(1:4,nrow=2,ncol=2,byrow=TRUE)
A
```

```
##      [,1] [,2]
## [1,]    1    2
## [2,]    3    4
```

Combining matrices

- Combine matrices with rbind() and cbind():

```
rbind(A,matrix(c(5,6),nrow=1,ncol=2))
```

```
##      [,1] [,2]
## [1,]    1    2
## [2,]    3    4
## [3,]    5    6
```

```
cbind(A,A)
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    2    1    2
## [2,]    3    4    3    4
```

Matrix attributes

- Matrices have a type, dimension (number of rows, number of columns) and optional attributes such as dimnames (row and column names).

```
typeof(A)
```

```
## [1] "integer"
```

```
dim(A)
```

```
## [1] 2 2
```

```r
colnames(A) <- c("var1","var2")
rownames(A) <- c("subj1","subj2")
A
```

```
##       var1 var2
## subj1    1    2
## subj2    3    4
```

```r
str(A)
```

```
##  int [1:2, 1:2] 1 3 2 4
##  - attr(*, "dimnames")=List of 2
##   ..$ : chr [1:2] "subj1" "subj2"
##   ..$ : chr [1:2] "var1" "var2"
```

Subsetting matrices

- Subset with [and a comma to separate rows from columns:

```
A[1,1]
```

```
## [1] 1
```

```
A[1,]
```

```
## var1 var2
##    1    2
```

```
A[,1]
```

```
## subj1 subj2
##     1     3
```

- When a subsetting operation leads to a vector, the dimension of the object is "dropped" from 2 to 1. To prevent this use drop=FALSE:

```
A[1,,drop=FALSE]
```

Extracting elements from matrices

▶ Can use [[to extract elements, but this is not necessary because of the way subsetting to a single element drops to a vector of length 1 by default:

```
A[[1,1]]
```

```
## [1] 1
```

```
A[1,1]
```

```
## [1] 1
```

Coercion: Matrices to/from vectors

- We have already seen how matrix() coerces a vector to a matrix
- as.vector() applied to a matrix creates a vector by concatenating columns:

```
as.vector(A)
```

```
## [1] 1 3 2 4
```

Data frames

- Data frames (class `data.frame`) are the usual way to store data in R.
    - Rows are intended to be observational units, columns variables
    - Implemented as a list (columns are list elements), but also behave like a matrix in terms of combining and subsetting.
- Create with `data.frame`:

```
set.seed(1)
n <- 4
x <- 1:n; y <- rnorm(n,mean=x,sd=1) # multiple commands separated by ;
dd <- data.frame(x=x,y=y) # like making a list
str(dd)
```

```
## 'data.frame':    4 obs. of  2 variables:
##  $ x: int  1 2 3 4
##  $ y: num  0.374 2.184 2.164 5.595
```

Subsetting and combining data frames like a list

```
dd$x
```

```
## [1] 1 2 3 4
```

```
dd[[1]]
```

```
## [1] 1 2 3 4
```

Subsetting and combining data frames like a matrix

```
dd[1:2,]
```

```
##   x         y
## 1 1 0.3735462
## 2 2 2.1836433
```

```
zz = data.frame(z=runif(4))
cbind(dd,zz)
```

```
##   x         y          z
## 1 1 0.3735462 0.62911404
## 2 2 2.1836433 0.06178627
## 3 3 2.1643714 0.20597457
## 4 4 5.5952808 0.17655675
```

Logical and relational operators

Logical operators

- The basic logical operators are described in `help("Logic")`.
- `!` is NOT
- `&` and `&&` are AND, with `&` acting vector-wise and `&&` acting on scalars
- `|` and `||` are OR, with `|` acting vector-wise and `||` acting on scalars
- Make sure you understand the following:

```r
x <- c(TRUE,TRUE,FALSE); y <- c(FALSE,TRUE,TRUE)
!x ; x&y ; x&&y ; x|y ; x||y
```

```
## [1] FALSE FALSE  TRUE
```

```
## [1] FALSE  TRUE FALSE
```

```
## [1] FALSE
```

```
## [1] TRUE TRUE TRUE
```

```
## [1] TRUE
```

- Notice how && and || act on the first element of the vectors x and y and ignore all the rest.

Relational operators

- ▶ Relational operators can be used to compare values in atomic vectors
  - ▶ See `help("Comparison")`
- ▶ `>` is greater than, `>=` is greater than or equal
- ▶ `<` is less than, `<=` is less than or equal
- ▶ `==` is equal and `!=` is not equal
- ▶ Make sure you understand the following:

```r
x <- 1:3; y <- 3:1
x>y ; x>=y ; x<y ; x<=y ; x==y ; x!=y
```

```
## [1] FALSE FALSE  TRUE

## [1] FALSE  TRUE  TRUE

## [1]  TRUE FALSE FALSE

## [1]  TRUE  TRUE FALSE

## [1] FALSE  TRUE FALSE
```

Subsetting vectors with logical expressions

- ▶ Can subset with logicals and [:

```
avec
```

```
##    age weight height
##     50    200     77
```

```
avec>100
```

```
##    age weight height
##  FALSE   TRUE  FALSE
```

```
avec[avec>100]
```

```
## weight
##    200
```

```
avec[avec>50 & avec<100]
```

```
## height
##     77
```

Subsetting matrices with logical expressions

- ▶ Can also subset matrices, but results may not be as expected:

```
A
```

```
##       var1 var2
## subj1    1    2
## subj2    3    4
```

```
A>1
```

```
##        var1 var2
## subj1 FALSE TRUE
## subj2  TRUE TRUE
```

```
A[A>1] # coerces to a vector
```

```
## [1] 3 2 4
```

Subset and assign with logical expressions

▶ Combine subset and assign to change the value of objects

```
A[A>1] <- 9
A
```

```
##       var1 var2
## subj1    1    9
## subj2    9    9
```

▶ In the above substitution, the vector 9 is shorter than the three elements in A>1 so R "recycles" the 9 three times.

Be careful about recycling:

```
A[A>1] <- c(-10,10) # Throws a warning
```

```
## Warning in A[A > 1] <- c(-10, 10): number of items to replace is not
## multiple of replacement length
```

```
A # R used c(-10,10), then just the -10
```

```
##       var1 var2
## subj1    1   10
## subj2  -10  -10
```

Aside: Special values

Missing values

- ▶ R has a special data code for missing data: `NA`
- ▶ Test for and set missing values with `is.na()`

```
avec
```

```
##    age weight height
##     50    200     77
```

```
is.na(avec)
```

```
##    age weight height
##  FALSE  FALSE  FALSE
```

```
is.na(avec) <- 2
avec
```

```
##    age weight height
##     50     NA     77
```

Infinite and undefined values

- R has a special codes for infinite values (`Inf`) and undefined values (`NaN`).
- Test for `Inf` and `NaN` with `is.infinite()` and `is.nan()`.

```
ii <- 1/0 ; nn <- 0/0
ii
```

```
## [1] Inf
```

```
is.infinite(ii)
```

```
## [1] TRUE
```

```
nn
```

```
## [1] NaN
```

```
is.nan(nn)
```

```
## [1] TRUE
```

The null object

- ▶ The null object, NULL, is an un-typed no-value object.
  - ▶ Test for NULL with is.null()
  - ▶ NULL can be used to initialize objects that will be created through combining, rbinding, etc.

```
x <- NULL; is.null(x)
```

```
## [1] TRUE
```

```
x <- c(x,1); x <- c(x,2); x
```

```
## [1] 1 2
```

```
# etc., or as a loop (more on these later)
x <- NULL
for(i in 1:2) {
  x <- c(x,i)
}
x
```

```
## [1] 1 2
```

R Functions

R functions

- Reference Wickham (2014), Advanced R, Chapter 6
- In R, functions are objects with three essential components:
    - the code inside the function, or `body`,
    - the list of arguments to the function, or `formals`, and
    - a data structure called an `environment` which is like a map to the memory locations of all objects defined in the function.
- Functions can have other attributes, but the above three are essential.

Example function

```
f <- function(x) {
  return(x^2)
}
f
```

```
## function(x) {
##   return(x^2)
## }
```

The function body

- This is the code we want to execute.
- When the end of a function is reached without a call to return(), the value of the last line is returned.
    - So in our example function, we could replace return(x^2) with just 'x^2.

The function formals

- ▶ These are the arguments to the function.
- ▶ Function arguments can have default values, as in:

```r
f <- function(x=0) { x^2}
```

- ▶ Argument defaults can be defined in terms of other arguments:

```r
f <- function(x=0,y=3*x) { x^2 + y^2 }
f()
```

```
## [1] 0
```

```r
f(x=1)
```

```
## [1] 10
```

```r
f(y=1)
```

```
## [1] 1
```

Argument matching when calling a function

- When you call a function, the arguments are matched first by name, then by "prefix" matching and finally by position:

```r
f <- function(firstarg,secondarg) {
  firstarg^2 + secondarg
}
f(firstarg=1,secondarg=2)
```

```
## [1] 3
```

```r
f(s=2,f=1)
```

```
## [1] 3
```

```r
f(2,f=1)
```

```
## [1] 3
```

```r
f(1,2)
```

```
## [1] 3
```

The function environment

- ▶ The environment within a function is like a map to the memory locations of all its variables.
- ▶ The function arguments are "passed by value", meaning that a copy is made and stored in the function's environment.
- ▶ Variables created within the function are also store in its environment

```r
f <- function(x) {
  y <- x^2
  ee <- environment() # Returns ID of environment w/in f
  print(ls(ee)) # list objects in ee
  ee
}
f(1) # function call
```

```
## [1] "ee" "x"  "y"
```

```
## <environment: 0x7fcddae1bc28>
```

Enclosing environments

- Our function f was defined in the global environment, .GlobalEnv, which "encloses" the environment within f.
- If f needs a variable and can't find it whithin f's environment, it will look for it in the enclosing environment, and then the enclosing environment of .GlobalEnv, and so on.
- The search() function lists the heirarchy of environments that enclose .GlobalEnv.

```
search()
```

```
## [1] ".GlobalEnv"        "package:stats"     "package:graphics
## [4] "package:grDevices" "package:utils"     "package:datasets
## [7] "package:methods"   "Autoloads"         "package:base"
```

- To facilitate this search, each environment includes a pointer to its enclosing environment.

R packages

Loading packages

- ► Use the `library()` command to load packages.
- ► When we load a package it is inserted in position 2 of the search list, just after `.GlobalEnv`.

```
# install.packages("hapassoc")
library(hapassoc)
search()
```

```
##  [1] ".GlobalEnv"        "package:hapassoc"  "package:stats"
##  [4] "package:graphics"  "package:grDevices" "package:utils"
##  [7] "package:datasets"  "package:methods"   "Autoloads"
## [10] "package:base"
```

Detaching packages

▶ Detach a package from the search list with `detach()`

```
detach("package:hapassoc")
search()
```

```
## [1] ".GlobalEnv"        "package:stats"     "package:graphics"
## [4] "package:grDevices" "package:utils"     "package:datasets"
## [7] "package:methods"   "Autoloads"         "package:base"
```

Package namespaces

- ▶ Package authors create a list of objects that will be visible to users when the package is loaded. This list is called the package namespace.
- ▶ You can access functions in a package's namespace without loading the package using the :: operator.

```
set.seed(321)
n<-30; x<-(1:n)/n; y<-rnorm(n,mean=x); ff<-lm(y~x)
car::sigmaHat(ff)
```

```
## [1] 0.926726
```

- ▶ Doing so does not add the package to the search list.