

# Lecture 7

Brad McNeney

2018-03-01

# Topics

- ▶ Tidying, or reshaping data
- ▶ Split-apply-combine

## Tidying (Reshaping) Data

# Tidying data with the tidyverse

```
library(tidyverse) # loads dplyr, ggplot2, tidyr, etc
```

- ▶ Reading: Hadley Wickham's article on tidy data, available from GitHub (Notes folder) or canvas.
  - ▶ Today's notes closely follow the `tidyr` vignette.
- ▶ “Tidy” data is ready for analysis, with one row for each sampled unit and columns for the variables measured on the units.
  - ▶ Often classify variables as “explanatory” or “response”
- ▶ Tabular data and repeated measures data are often not in tidy form.
- ▶ Examples:
  - ▶ Tabular data on new tuberculosis cases from WHO
  - ▶ Repeated measures from Billboard Top 100

# Tuberculosis (TB) cases

```
tb <- read.csv("tb.csv",stringsAsFactors=FALSE)
dim(tb)
```

```
## [1] 5769    23
```

```
tb[1:4,1:6]
```

```
##   iso2 year new_sp new_sp_m04 new_sp_m514 new_sp_m014
## 1   AD 1989     NA          NA          NA          NA
## 2   AD 1990     NA          NA          NA          NA
## 3   AD 1991     NA          NA          NA          NA
## 4   AD 1992     NA          NA          NA          NA
```

```
names(tb)[1:20]
```

```
## [1] "iso2"          "year"          "new_sp"        "new_sp_m04"
## [5] "new_sp_m514"   "new_sp_m014"   "new_sp_m1524"  "new_sp_m2534"
## [9] "new_sp_m3544"  "new_sp_m4554"  "new_sp_m5564"  "new_sp_m65"
## [13] "new_sp_mu"     "new_sp_f04"    "new_sp_f514"   "new_sp_f014"
## [17] "new_sp_f1524"  "new_sp_f2534"  "new_sp_f3544"  "new_sp_f4554"
```

# Structure of TB table

- ▶ First column is 2-letter country code, second is year, third is number of new cases for that country/year.
- ▶ Then come TB counts for different gender/age categories.
  - ▶ `new_sp` is “new cases by positive pulmonary smear assay”
  - ▶ gender is `m` or `f`
  - ▶ two special age categories 0-4, 5-14,
  - ▶ age categories 0-14, 15-24, 25-34, 35-44, 45-54, 55-65, 65+, unknown (`u`)
- ▶ Gender/age columns are not variables, they are data on the observed units.
- ▶ Tidy data would have one row for each country, year, gender and age category, with a column of counts

# Billboard Top 100 rankings of songs

```
bb <- read.csv("billboard.csv",stringsAsFactors = FALSE)
dim(bb)
```

```
## [1] 317 83
```

```
bb[1:3,1:6]
```

```
##   year artist.inverted      track time genre date.entered
## 1 2000 Destiny's Child Independent Women Part I 3:38  Rock    2000-09-23
## 2 2000           Santana             Maria, Maria 4:18  Rock    2000-02-12
## 3 2000   Savage Garden      I Knew I Loved You 4:07  Rock    1999-10-23
```

```
names(bb)[c(1:10,ncol(bb))]
```

```
## [1] "year"           "artist.inverted" "track"
## [4] "time"           "genre"           "date.entered"
## [7] "date.peaked"    "x1st.week"       "x2nd.week"
## [10] "x3rd.week"      "x76th.week"
```

# Structure of the Billboard table

- ▶ Columns `year` through `date.peaked` describe the song, then `x1st.week` through `x76th.week` are the chart positions for the first through 76th weeks.
  - ▶ If a song is on the chart for less than 76 weeks, its position is NA for any missing weeks.
- ▶ Weeks are not variables, they are the time data for the time series.



# Tidying the Billboard data

- ▶ Main step is to consolidate, or “gather” the rankings in the different weeks into a `rank` variable.
- ▶ Before gathering, will select/rename some of the variables.
- ▶ After gathering, will create some new variables and sort the data frame.

# Select and rename

- ▶ Won't need time or genre.
  - ▶ `select()` from `dplyr` can use `-` to de-select
- ▶ Rename `artist.inverted`
  - ▶ `rename()` from `dplyr` takes arguments of the form `newname = oldname`

```
bb <-  
bb %>% select(-time, -genre) %>%  
rename(artist = artist.inverted)
```

## Gather the weeks into a “long” version of the Billboard data

- ▶ Leave each song info variable as-is.
- ▶ The data, or “values”, are the chart positions.
- ▶ The weeks are descriptors or “keys” for these values.
- ▶ We want to create key-value pairs for each observation.
  - ▶ There will be missing values, which we can remove.
- ▶ The `gather()` function from `tidyr` gathers specified columns into keys (e.g., `week`) and values (e.g., `rank`).

## gather() for the Billboard data

```
# gather (data, key, value, ... ) where ... are the columns to collapse  
bblong <- gather(bb, week, rank, x1st.week:x76th.week, na.rm=TRUE)  
head(bblong, n=4)
```

##	year	artist	track	date.entered	date.peaked
## 1	2000	Destiny's Child	Independent Women Part I	2000-09-23	2000-11-18
## 2	2000	Santana	Maria, Maria	2000-02-12	2000-04-08
## 3	2000	Savage Garden	I Knew I Loved You	1999-10-23	2000-01-29
## 4	2000	Madonna	Music	2000-08-12	2000-09-16

  

##	week	rank
## 1	x1st.week	78
## 2	x1st.week	15
## 3	x1st.week	71
## 4	x1st.week	41

## More cleaning suggested in the vignette

- ▶ Extract week numbers from week variable
- ▶ Coerce date.entered to a Date object
- ▶ Calculate the date of each ranking based on the date it entered the charts and the week.
- ▶ Sort (“arrange”) on artist, track and week.

```
bb <-  
  bblong %>% mutate(week = parse_number(week),  
                    date = as.Date(date.entered) + 7*(week-1)) %>%  
  select(-date.entered) %>% # don't need date.entered anymore  
  arrange(artist,track,week)  
head(bb,n=3)
```

##	year	artist	track	date.peaked	week	rank
## 1	2000	2 Pac Baby Don't Cry (Keep Ya Head Up II)	2000-03-11	1	87	
## 2	2000	2 Pac Baby Don't Cry (Keep Ya Head Up II)	2000-03-11	2	82	
## 3	2000	2 Pac Baby Don't Cry (Keep Ya Head Up II)	2000-03-11	3	72	
##		date				
## 1	2000-02-26					
## 2	2000-03-04					
## 3	2000-03-11					

# Tidying the TB data

- ▶ Recall structure of the data: country, year, count of new cases, counts of new cases by gender/age categories.

```
names(tb)[1:10]
```

```
## [1] "iso2"          "year"          "new_sp"        "new_sp_m04"  
## [5] "new_sp_m514"   "new_sp_m014"   "new_sp_m1524"  "new_sp_m2534"  
## [9] "new_sp_m3544"  "new_sp_m4554"
```

- ▶ Main step is to “gather” TB prevalence in the different gender/age categories into a count variable.
  - ▶ Complicated by the coding of gender/age categories
- ▶ Before gathering, will remove unneeded variables and add country names to supplement 2-letter codes.

# Remove variables

- ▶ Won't need overall count
- ▶ Special categories 0-4 and 5-14 overlap with 0-14 so remove
- ▶ Age unknown not useful for analysing trends, so remove

```
tb <- select(tb, -new_sp, -contains("04"), -contains("514"),  
             -new_sp_mu, -new_sp_fu)  
tb[1:3, 1:10]
```

```
##   iso2 year new_sp_m014 new_sp_m1524 new_sp_m2534 new_sp_m3544  
## 1   AD 1989           NA             NA             NA             NA  
## 2   AD 1990           NA             NA             NA             NA  
## 3   AD 1991           NA             NA             NA             NA  
## new_sp_m4554 new_sp_m5564 new_sp_m65 new_sp_f014  
## 1           NA           NA           NA           NA  
## 2           NA           NA           NA           NA  
## 3           NA           NA           NA           NA
```

# Add country names to supplement country codes

- ▶ I found a translation of the ISO-2 country codes at [<http://data.okfn.org/data/core/country-list>] and saved as `countryCodes.csv` in the Notes folder.

```
cc <- read.csv("countryCodes.csv", stringsAsFactors = FALSE)
# cc has columns "Name" and "Code". "Code" matches "iso2" in tb
```



- Exercise: Find out which ISO-2 codes are in `tb` but not in `countryCodes.csv`, google the missing codes, and add the country names to `cc` manually.

```
unique(tb$iso2[!(tb$iso2 %in% cc$Code)])
```

```
## [1] "AN" "YU"
```

```
cc <- rbind(cc, data.frame(Name=c("Netherlands Antilles", "Yugoslavia"),  
                           Code=c("AN", "YU")))  
tb <- inner_join(cc, tb, by = c("Code" = "iso2"))  
tb[1:2, 1:6]
```

```
##           Name Code year new_sp_m014 new_sp_m1524 new_sp_m2534  
## 1 Afghanistan  AF 1980           NA           NA           NA  
## 2 Afghanistan  AF 1981           NA           NA           NA
```

## Gather counts for demographic groups

- Create demographic variable `demog` and count variable `count` by gathering over all variables except `Name`, `Code` and `year`.

```
tblong <- gather(tb,demog,count,-Name,-Code,-year,na.rm=TRUE)
head(tblong)
```

##	Name	Code	year	demog	count
## 13	Afghanistan	AF	1997	new_sp_m014	0
## 14	Afghanistan	AF	1998	new_sp_m014	30
## 15	Afghanistan	AF	1999	new_sp_m014	8
## 16	Afghanistan	AF	2000	new_sp_m014	52
## 17	Afghanistan	AF	2001	new_sp_m014	129
## 18	Afghanistan	AF	2002	new_sp_m014	90

## Separate gender from age category.

- First remove new\_sp\_, then separate result on first column (help(separate))

```
maxlen <- max(nchar(tblong$demog))
tblong %>% mutate(demog = substr(demog,8,maxlen)) %>%
  separate(demog, into=c("gender","agecat"),sep=1) -> tb
head(tb)
```

##	Name	Code	year	gender	agecat	count
## 1	Afghanistan	AF	1997	m	014	0
## 2	Afghanistan	AF	1998	m	014	30
## 3	Afghanistan	AF	1999	m	014	8
## 4	Afghanistan	AF	2000	m	014	52
## 5	Afghanistan	AF	2001	m	014	129
## 6	Afghanistan	AF	2002	m	014	90

```
save(tb,file="tb.RData")
```

Split-Apply-Combine

## Subgroup summaries

- ▶ Data visualization and modelling is often in terms of subgroups.
- ▶ Illustrate with some data on enrollments in Stat and Act Sci courses over the 2007/08 to 2015/16 academic years.
  - ▶ Data on full-time equivalents (FTEs, equal to 30 credit hours taught) by year and course.
- ▶ Recurring theme: Need to split the data into subgroups, transform or summarize, and reassemble, or unsplit.
  - ▶ Has come to be known as “split-apply-combine”

# Science enrollments database

- ▶ Load the scilong data frame created by FTE.Rmd
  - ▶ Look through the FTE.Rmd script if you haven't already.

```
library(tidyverse)
load("scilong.RData")
head(scilong)
```

##	Subject	CrsNum	CreditHrs	semester	enrollment	FTEs	year
## 1	ACMA	210	3	1077	51	5.10000	2008
## 2	ACMA	335	3	1077	20	2.00000	2008
## 3	ACMA	425	3	1077	22	2.20000	2008
## 4	ACMA	465	3	1077	16	1.60000	2008
## 5	ACMA	490	3	1077	4	0.40000	2008
## 6	BISC	100	4	1077	176	23.46667	2008

# Stat and Act Sci data

```
stat <- filter(scilong,Subject=="STAT" | Subject=="ACMA")
head(stat)
```

##	Subject	CrsNum	CreditHrs	semester	enrollment	FTEs	year
## 1	ACMA	210	3	1077	51	5.1	2008
## 2	ACMA	335	3	1077	20	2.0	2008
## 3	ACMA	425	3	1077	22	2.2	2008
## 4	ACMA	465	3	1077	16	1.6	2008
## 5	ACMA	490	3	1077	4	0.4	2008
## 6	STAT	100	3	1077	49	4.9	2008

## Split-apply-combine example 1: yearly percent FTEs

- ▶ Suppose we want the percent of FTEs in a year that are attributable to each course taught.
- ▶ Split the data by year, compute proportion of FTEs for each course in that year, and combine the proportions into a variable that can be included in the stat data frame.
- ▶ Illustrate base R and `dplyr` approaches.



## Example 1: split

- ▶ The base R function `split()` splits a data frame on a grouping variable, which is a vector or list of vectors that can be coerced to `factor(s)`, and returns a list.

```
sp.stat <- split(stat, stat$year)
names(sp.stat)
```

```
## [1] "2008" "2009" "2010" "2011" "2012" "2013" "2014" "2015" "2016"
```

```
head(sp.stat[["2008"]])
```

```
##   Subject CrsNum CreditHrs semester enrollment FTEs year
## 1    ACMA   210         3      1077         51  5.1 2008
## 2    ACMA   335         3      1077         20  2.0 2008
## 3    ACMA   425         3      1077         22  2.2 2008
## 4    ACMA   465         3      1077         16  1.6 2008
## 5    ACMA   490         3      1077          4  0.4 2008
## 6    STAT   100         3      1077         49  4.9 2008
```

```
str(sp.stat[["2008"]])
```

```
## 'data.frame':   47 obs. of  7 variables:
## $ Subject      : chr  "ACMA" "ACMA" "ACMA" "ACMA" ...
```

## Example 1: Split, cont.

```
sp.stat <- split(stat,list(stat$year,stat$Subject))  
names(sp.stat)
```

```
## [1] "2008.ACMA" "2009.ACMA" "2010.ACMA" "2011.ACMA" "2012.ACMA"  
## [6] "2013.ACMA" "2014.ACMA" "2015.ACMA" "2016.ACMA" "2008.STAT"  
## [11] "2009.STAT" "2010.STAT" "2011.STAT" "2012.STAT" "2013.STAT"  
## [16] "2014.STAT" "2015.STAT" "2016.STAT"
```

```
head(sp.stat[["2008.STAT"]])
```

##	Subject	CrsNum	CreditHrs	semester	enrollment	FTEs	year
## 6	STAT	100	3	1077	49	4.9	2008
## 7	STAT	101	3	1077	59	5.9	2008
## 8	STAT	201	3	1077	284	28.4	2008
## 9	STAT	203	3	1077	164	16.4	2008
## 10	STAT	270	3	1077	185	18.5	2008
## 11	STAT	285	3	1077	47	4.7	2008

## group\_by() from dplyr

- ▶ Call is similar to `split`, but we specify multiple variables to group on by comma-separated names.
- ▶ Output is a tibble (data frame with some different default behaviours).

```
sp.stat.dplyr <- group_by(stat,year,Subject)
sp.stat.dplyr
```

```
## # A tibble: 468 x 7
## # Groups:   year, Subject [18]
##   Subject CrsNum CreditHrs semester enrollment   FTEs   year
##   <chr>    <chr>      <int>    <dbl>      <int> <dbl> <dbl>
## 1 ACMA     210           3    1077.        51  5.10  2008.
## 2 ACMA     335           3    1077.        20  2.00  2008.
## 3 ACMA     425           3    1077.        22  2.20  2008.
## 4 ACMA     465           3    1077.        16  1.60  2008.
## 5 ACMA     490           3    1077.         4  0.400 2008.
## 6 STAT     100           3    1077.        49  4.90  2008.
## 7 STAT     101           3    1077.        59  5.90  2008.
## 8 STAT     201           3    1077.       284 28.4   2008.
## 9 STAT     203           3    1077.       164 16.4   2008.
## 10 STAT    270           3    1077.       185 18.5   2008.
## # ... with 458 more rows
```

## Example 1: Apply

- ▶ Create a new variable `FTEproportion = FTEs/sum(FTEs)` for each sub-group data frame and save the new variable in the respective data frames.
- ▶ Can use the base R function `lapply()`
  - ▶ stands for “list apply” – apply a function to each element of a list and return a list as output
- ▶ It turns out the following call to `lapply()` does what we want.

```
tem <- lapply(sp.stat,transform,FTEproportion=FTEs/sum(FTEs))
```

- ▶ To see why, start with simpler uses of `lapply()`.

## Simpler example of lapply()

- Define a function to apply to each list element and apply it:

```
fsum <- function(x) { # x is a list element  
  sum(x$FTEs) # assumes list elements have an FTEs column  
}  
tem <- lapply(sp.stat,fsum)  
tem[1:2]
```

```
## $`2008.ACMA`  
## [1] 20.36667  
##  
## $`2009.ACMA`  
## [1] 18.63333
```

## Simpler example, cont.

- If our function takes more arguments than just the list element, we add them after the function name.

```
fsum <- function(x,cname) {  
  sum(x[,cname])  
}  
tem <- lapply(sp.stat,fsum,"FTEs")  
tem[1:2]
```

```
## $`2008.ACMA`  
## [1] 20.36667  
##  
## $`2009.ACMA`  
## [1] 18.63333
```

## Our use of `lapply()`

- ▶ Adding a column to each sub-group data frame requires a function that takes the data frame as an argument and returns the augmented version.
  - ▶ This is what `transform()` does

```
head(transform(sp.stat[[1]], FTEproportion = FTEs/sum(FTEs)))
```

##	Subject	CrsNum	CreditHrs	semester	enrollment	FTEs	year	FTEproportion
## 1	ACMA	210	3	1077	51	5.1	2008	0.25040917
## 2	ACMA	335	3	1077	20	2.0	2008	0.09819967
## 3	ACMA	425	3	1077	22	2.2	2008	0.10801964
## 4	ACMA	465	3	1077	16	1.6	2008	0.07855974
## 5	ACMA	490	3	1077	4	0.4	2008	0.01963993
## 19	ACMA	315	3	1081	21	2.1	2008	0.10310966

## Putting it all together

```
sp.stat <- lapply(sp.stat,transform,FTEproportion=FTEs/sum(FTEs))  
head(sp.stat[[1]])
```

##	Subject	CrsNum	CreditHrs	semester	enrollment	FTEs	year	FTEproportion
## 1	ACMA	210	3	1077	51	5.1	2008	0.25040917
## 2	ACMA	335	3	1077	20	2.0	2008	0.09819967
## 3	ACMA	425	3	1077	22	2.2	2008	0.10801964
## 4	ACMA	465	3	1077	16	1.6	2008	0.07855974
## 5	ACMA	490	3	1077	4	0.4	2008	0.01963993
## 19	ACMA	315	3	1081	21	2.1	2008	0.10310966



## Detour: The apply family of functions in R

- The “original” apply is `apply()`, which can be used to apply a function to rows or columns of a matrix.

```
mat <- matrix(1:6,ncol=2,nrow=3)
mat
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
apply(mat,1,sum) # row-wise sums; rowSums() is faster
```

```
## [1] 5 7 9
```

```
apply(mat,2,sum) # column-wise; colSums() is faster
```

```
## [1] 6 15
```

## Detour, cont.

- ▶ `sapply()` takes the output of `lapply()` and simplifies to a vector or matrix.

```
fsum <- function(x) { sum(x$FTEs) }  
sapply(sp.stat,fsum)[1:2]
```

```
## 2008.ACMA 2009.ACMA  
## 20.36667 18.63333
```

## Detour, cont.

- ▶ Other apply-like functions `vapply()`, `mapply()`, `tapply()`, ...
- ▶ I don't use these.
  - ▶ See their respective help pages for information.

# The apply step with dplyr

- ▶ Actions (“verbs”) like `mutate()` are applied to the data within groups when passed a grouped object.
  - ▶ That is, the data table is broken into groups and `mutate()` is applied separately to each group.

```
sp.stat.dplyr <- mutate(sp.stat.dplyr, FTEpp = FTEs/sum(FTEs))  
select(sp.stat.dplyr, Subject, FTEs, year, FTEpp)
```

```
## # A tibble: 468 x 4  
## # Groups:   year, Subject [18]  
##   Subject    FTEs  year  FTEpp  
##   <chr>    <dbl> <dbl> <dbl>  
## 1 ACMA      5.10  2008.  0.250  
## 2 ACMA      2.00  2008.  0.0982  
## 3 ACMA      2.20  2008.  0.108  
## 4 ACMA      1.60  2008.  0.0786  
## 5 ACMA      0.400 2008.  0.0196  
## 6 STAT      4.90  2008.  0.0208  
## 7 STAT      5.90  2008.  0.0250  
## 8 STAT     28.4   2008.  0.121  
## 9 STAT     16.4   2008.  0.0696  
## 10 STAT     18.5   2008.  0.0785  
## # ... with 458 more rows
```

## The combine step

- ▶ The base R function `unsplit()` will combine the elements of the list that was generated by `split()`
- ▶ Pass `unsplit()` the list of variables used to define the splits.

```
head(unsplit(sp.stat,list(stat$year,stat$Subject)))
```

##	Subject	CrsNum	CreditHrs	semester	enrollment	FTEs	year	FTEproportion
## 1	ACMA	210	3	1077	51	5.1	2008	0.25040917
## 2	ACMA	335	3	1077	20	2.0	2008	0.09819967
## 3	ACMA	425	3	1077	22	2.2	2008	0.10801964
## 4	ACMA	465	3	1077	16	1.6	2008	0.07855974
## 5	ACMA	490	3	1077	4	0.4	2008	0.01963993
## 6	STAT	100	3	1077	49	4.9	2008	0.02079796

# The combine step with dplyr

- Use `ungroup()`

```
ungroup(sp.stat.dplyr)
```

```
## # A tibble: 468 x 8
```

```
##   Subject CrsNum CreditHrs semester enrollment   FTEs   year  FTEpp
##   <chr>   <chr>    <int>    <dbl>    <int>  <dbl>  <dbl>  <dbl>
## 1 ACMA    210         3    1077.     51  5.10  2008.  0.250
## 2 ACMA    335         3    1077.     20  2.00  2008.  0.0982
## 3 ACMA    425         3    1077.     22  2.20  2008.  0.108
## 4 ACMA    465         3    1077.     16  1.60  2008.  0.0786
## 5 ACMA    490         3    1077.      4  0.400 2008.  0.0196
## 6 STAT    100         3    1077.     49  4.90  2008.  0.0208
## 7 STAT    101         3    1077.     59  5.90  2008.  0.0250
## 8 STAT    201         3    1077.    284 28.4   2008.  0.121
## 9 STAT    203         3    1077.    164 16.4   2008.  0.0696
## 10 STAT   270         3    1077.    185 18.5   2008.  0.0785
## # ... with 458 more rows
```

# Summary of split-apply-combine

## ► Base R:

```
sp.stat <- split(stat,list(stat$year,stat$Subject))  
sp.stat <- lapply(sp.stat,transform,FTEproportion = FTEs/sum(FTEs))  
stat <- unsplit(sp.stat,list(stat$year,stat$Subject))
```

## ► dplyr

```
stat %>% group_by(year,Subject) %>%  
  mutate(FTEproportion = FTEs/sum(FTEs)) %>%  
  ungroup() -> stat  
save(stat,file="statEnrol.RData")
```

## Split-apply-combine with summarise()

- ▶ In the apply step, we may wish to calculate some sort of summary, rather than a transformation of a variable.
- ▶ For example, suppose we want to calculate total FTEs by year and subject, and return a data frame

```
stat %>% group_by(year,Subject) %>%  
  summarise(totalFTEs = sum(FTEs)) %>%  
  ungroup() -> totals  
head(totals,n=4)
```

```
## # A tibble: 4 x 3  
##   year Subject totalFTEs  
##   <dbl> <chr>      <dbl>  
## 1 2008. ACMA        20.4  
## 2 2008. STAT       236.  
## 3 2009. ACMA        18.6  
## 4 2009. STAT       213.
```



# Split-apply-combine with `lapply()`

- Compare to base R

```
tem <- split(stat,list(stat$year,stat$Subject))  
tem <- lapply(tem,function(x) sum(x$FTEs))  
tem[1:4]
```

```
## $`2008.ACMA`  
## [1] 20.36667  
##  
## $`2009.ACMA`  
## [1] 18.63333  
##  
## $`2010.ACMA`  
## [1] 23.06667  
##  
## $`2011.ACMA`  
## [1] 24.03333
```

- Then would have to write code to coerce output to a data frame.